

# Discretization-Agnostic Deep Self-Supervised 3D Surface Parameterization

Chandradeep Pokhariya \*

Shanthika Naik\*

IIIT Hyderabad

India

Astitva Srivastava

Avinash Sharma

IIIT Hyderabad

India

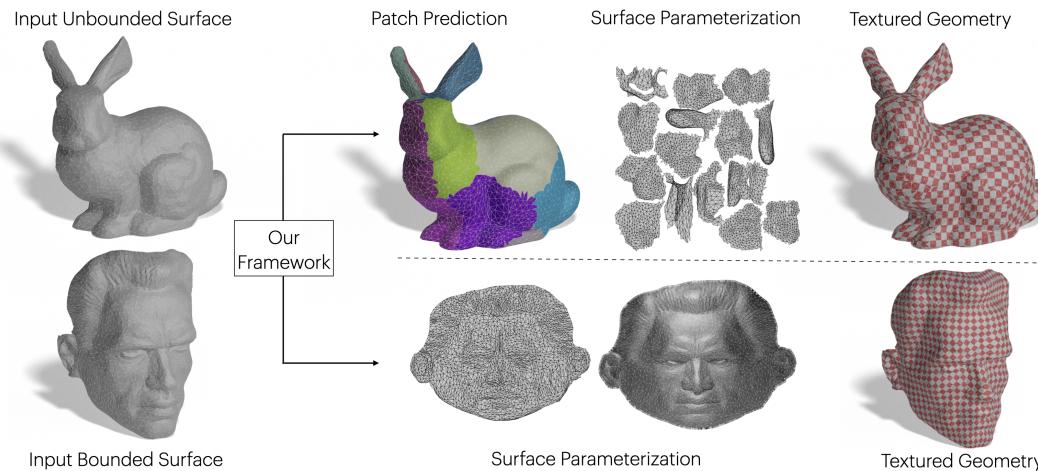


Figure 1: UV parameterization for bounded and unbounded surfaces estimated via our proposed framework.

## ABSTRACT

We present a novel self-supervised framework for learning the discretization-agnostic surface parameterization of arbitrary 3D objects with both bounded and unbounded surfaces. Our framework leverages diffusion-enabled global-to-local shape context for each vertex first to partition the unbounded surface into multiple patches using the proposed self-supervised PatchNet and subsequently perform independent UV parameterization of these patches by learning forward and backward UV mapping for individual patches. Thus, our framework enables learning a discretization-agnostic parameterization at a lower resolution and then directly inferring the parameterization for a higher-resolution mesh without retraining. We evaluate our framework on multiple 3D objects from the publicly available SHREC [Lian et al. 2011] dataset and report superior/faster UV parameterization over conventional methods.

\*Both authors contributed equally to this research.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SA '22 Technical Communications, December 6–9, 2022, Daegu, Republic of Korea  
© 2022 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 978-1-4503-9465-9/22/12...\$15.00  
<https://doi.org/10.1145/3550340.3564235>

## CCS CONCEPTS

- Computing methodologies → Parametric curve and surface models; Neural networks.

## KEYWORDS

UV parameterization, texture mapping, neural network, self-supervised learning, surface parameterization.

## ACM Reference Format:

Chandradeep Pokhariya, Shanthika Naik, Astitva Srivastava, and Avinash Sharma. 2022. Discretization-Agnostic Deep Self-Supervised 3D Surface Parameterization. In *SIGGRAPH Asia 2022 Technical Communications (SA '22 Technical Communications)*, December 6–9, 2022, Daegu, Republic of Korea. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3550340.3564235>

## 1 INTRODUCTION

Estimating the UV parameterization of arbitrary 3D surfaces lies at the core of computer graphics and geometry processing domain, with a wide range of applications such as 3D modelling, texture-mapping, remeshing, simulation, etc. Formally, it is defined as the projection of vertices of a tessellated surface (polygon mesh) onto a 2D map (UV plane). Determining the aforementioned mapping is not a trivial task and demands a solution with specific properties. The estimated mapping is expected to be isometric, conformal, and non-overlapping. Existing conventional methods [Lévy et al. 2002; Li et al. 2018; Sander et al. 2001; Sawhney and Crane 2017; Wang et al. 2013] aim to estimate an object-centric mapping with an iterative optimization process, focusing on minimizing an energy function explicitly constructed to retain the desired properties. However,

they face scalability issues while dealing with high-resolution object meshes and are also prone to local minima.

With the advent of deep learning, researchers are harnessing the power of neural networks to solve various ill-posed problems, offering tractable solutions. Neural surface parameterization has recently been attempted [Aigerman et al. 2022] but under supervised, data-driven settings, requiring a large amount of training data. Such supervised learning solutions get subjected to data bias and hence suffer from poor generalization to unseen, out-of-distribution samples.

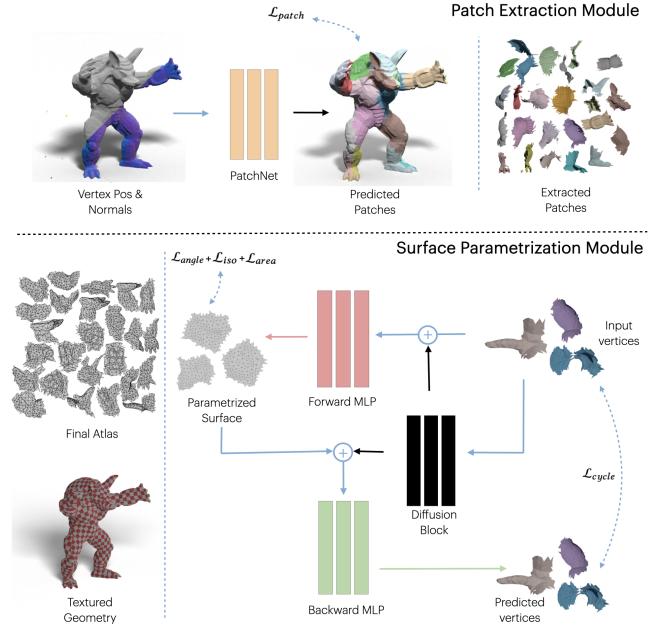
This paper presents a novel, self-supervised framework for learning the discretization-agnostic surface parameterization of arbitrary 3D objects with both bounded and unbounded surfaces as shown in Figure 1. First, to handle unbounded surfaces (e.g., a sphere) or surfaces with regions of extreme extrinsic curvature, we propose a learning-based partitioning of the given surface into multiple bounded patches, which are independently parameterized. To this end, we employ a self-supervised network that assigns each 3D point of the surface to one of the patches, trained using losses based on local features (such as face-normals) and geodesic relationships within the patch.

Subsequently, we propose to learn the surface parameterization of an arbitrary (bounded) 3D surface to a UV plane using a *Multi-layer Perceptron* (MLP). More specifically, given a *bounded* 3D surface (patch), we train the forward MLP to predict per-point UV coordinates independently. In order to ensure a meaningful UV mapping, we enforce cycle-consistency loss between the input and reconstructed surface by learning a backward mapping (UV-to-3D) MLP. Additional losses are employed to achieve desired properties of surface parameterization, i.e., isometric, conformal, and area-preserving. A diffusion process [Sharp et al. 2020] over the mesh provides a multi-scale characterization of the underlying surface, entailing a global-to-local context for each vertex. Hence, the DiffusionNet backbone is used for PatchNet, and similarly, respective features are appended while learning surface parameterization to achieve discretization-agnostic UV mapping. A key advantage of learning a discretization-agnostic parameterization is that we can learn on meshes at a lower resolution and then directly infer the parameterization for high resolution meshes without retraining, as shown in Figure 3.

## 2 RELATED WORK

Conventional methods for mesh parameterization can be categorised as *single-patch, fixed boundary* methods, e.g., harmonic parameterization [Wang et al. 2013], *single-patch, free boundary* methods, e.g., LSCM [Lévy et al. 2002]; and *global parameterization* methods, e.g., Voronoi atlas parameterization [Sander et al. 2001]. Global parameterization methods can deal with meshes of arbitrary genus. They achieve this by partitioning the unbounded mesh into multiple bounded patches or detecting one or more seams to cut the mesh, making it bounded. Boundary-First Flattening [Sawhney and Crane 2017] and OptCuts [Li et al. 2018] fall into this class.

Neural parameterization methods have been gaining popularity over the past few years due to their ability to address ill-posed problems. AtlasNet [Groueix et al. 2018], and DGP [Williams et al. 2018] propose a way of surface reconstruction and parameterization by training a neural network to represent a single UV chart



**Figure 2: The outline of proposed framework.**

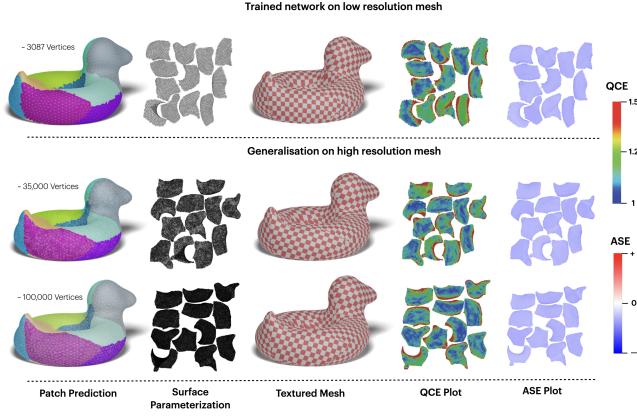
over the reconstructed surface. Both methods use a fixed number of patches for the surface parameterization but require a different neural network for every patch, which is overkill and difficult to scale. A recent work AUV-Net [Chen et al. 2022] takes a point cloud as input and learns parameterization of aligned surfaces (e.g., faces and humans in T-poses) using a cycle-loss, but requires all the meshes to have similar topology and same orientation to enable learning. Moreover, the proposed two-patch estimation method is very naive and cannot scale to an arbitrary number of patches. Another recent method [Aigerman et al. 2022] learns intrinsic mapping of arbitrary surfaces in a supervised setup, where a conventional method acts as the ground truth.

## 3 METHOD

We now describe the proposed framework in detail. The input to our framework is a mesh  $\mathcal{M} = \{\mathcal{V}, \mathcal{F}, \mathcal{N}_{\mathcal{V}}\}$ , where  $\mathcal{V}$ ,  $\mathcal{F}$  and  $\mathcal{N}_{\mathcal{V}}$  are the sets of vertex positions, faces and vertex-normals respectively. Our framework consists of two modules: (i) Patch extraction module and (ii) Surface parameterization module.

### 3.1 Patch Extraction Module

Handling surfaces with regions of high extrinsic curvature or unbounded topology requires the 3D manifold to be partitioned into multiple bounded patches to minimize distortion and overlap. Each patch is defined as  $\mathcal{P}_k = \{\mathcal{V}_k, \mathcal{F}_k, \mathcal{N}_{F_k}\}$  ( $k = 1, 2, \dots, K$ ), where  $\mathcal{V}_k \subseteq \mathcal{V}$  is the set of vertices belonging to  $\mathcal{P}_k$ .  $\mathcal{F}_k \subseteq \mathcal{F}$  is the set of faces defined on  $\mathcal{V}_k$  and  $\mathcal{N}_{F_k} \subseteq \mathcal{N}_{\mathcal{F}}$  is the associated set of face-normals. We propose PatchNet with parameters  $\phi_{patch}$ , which learns to assign each vertex of  $\mathcal{M}$  to one of the  $K$  patches, as shown in Figure 2. Here,  $K$  is a controllable parameter and can vary based on the acceptable amount of distortion in the input mesh. To learn



**Figure 3: Discretization-agnostic UV parameterization.**

the parameters  $\phi_{patch}$ , we minimize the following cosine similarity constraint on the estimated patches:

$$\mathcal{L}_{cos} = \sum_{k=1}^K \frac{1}{|\mathcal{F}_k|} \left[ 1 - \left( \sum_{i,j \in \mathcal{F}_k} (\hat{n}_i^T \hat{n}_j) \right) \right]^2 \quad (1)$$

where  $i, j \in \mathcal{F}_k$  are the pair of faces with unit normal vectors  $\hat{n}_i, \hat{n}_j \in \mathcal{N}_{F_k}$ , respectively, and  $|\mathcal{F}_k|$  is the number of faces in that patch. The above constraint has the effect of producing locally flat patches. However, geodesically far-apart triangles with high cosine similarity may be assigned to the same patch, which is undesirable. To circumvent such disjoint assignments, we minimize the following additional constraint:

$$\mathcal{L}_{geo} = \sum_{k=1}^K \frac{1}{|\mathcal{P}_k|} \left( \sum_{i,j \in \mathcal{V}_k} g(i, j) \right) \quad (2)$$

where  $g(i, j)$  denotes the geodesic distance between the pair of vertices  $i$  &  $j$  within the patch and  $|\mathcal{P}_k|$  is the number of vertices in that patch. We model PatchNet using DiffusionNet [Sharp et al. 2020] architecture to achieve multi-scale characterization of the underlying surface, entailing a global-to-local context for all the vertices. Input to PatchNet is the vertices  $\mathcal{V}$  and vertex-normals  $\mathcal{N}_V$ , and the output is the predicted assignment probability for all the vertices to each of the  $K$  patches. Subsequently, per-face probabilities are obtained by taking the mean probabilities of the corresponding face vertices. We further consolidate the per-face probabilities by taking an average over neighboring faces, and then each face is assigned to the patch with the highest probability. Note that the whole mesh can be considered as a single patch in the case of a bounded surface with extrinsic curvature of low variability. The combined objective function for patch extraction becomes  $\mathcal{L}_{patch} = \lambda_{cos} \mathcal{L}_{cos} + \lambda_{geo} \mathcal{L}_{geo}$ .

### 3.2 Surface Parameterization Module

Each patch  $\mathcal{P}_k = \{\mathcal{V}_k, \mathcal{F}_k, \mathcal{N}_k\}$  is treated as a separate bounded surface and is independently parameterized. Let  $f : \mathbb{R}^3 \rightarrow \mathbb{R}^2$  be the mapping of each vertex  $v \in \mathcal{V}_k$  to a 2D point  $u$  on the UV plane. We propose to represent  $f$  using a *forward* mapping network  $MLP_f$

with learnable parameters  $\phi_f$ . First, the set of vertices  $\mathcal{V}_k$  for the given patch is passed to the diffusion block to get a global shape encoding  $\psi \in \mathbb{R}^{128}$ . Per-vertex input given to  $MLP_f$  is  $z \in \mathbb{R}^{131}$  ( $v$  concatenated with  $\psi$ ) and the output is  $u \in \mathbb{R}^2$  (UV coordinate), i.e.  $u = MLP_f(z)$ . Since we do not have corresponding ground truth UV coordinates, we resort to a self-supervised cycle-consistency loss. We employ another  $MLP_{f^{-1}}$  with learnable parameters  $\phi_{f^{-1}}$  to represent the *backward* mapping  $f^{-1} : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ .  $MLP_{f^{-1}}$  takes  $u$  as input and predicts its corresponding 3D position, which ideally should match with the input vertex position  $v$ . We enforce this consistency by minimizing the following cycle loss:

$$\mathcal{L}_{cycle} = \frac{1}{|\mathcal{V}_k|} \sum_{v \in \mathcal{V}_k} (v - MLP_{f^{-1}}(u))^2 \quad (3)$$

Note that due to presence of non-linear activation functions in  $MLP_f$  and  $MLP_{f^{-1}}$ , the condition  $\phi_f \cdot \phi_{f^{-1}} = I$  need not hold. Per-vertex prediction can be noisy, resulting in an irregular UV space. Conditioning the MLPs with the diffusion-based global shape-encoding  $\psi$  regularizes the UV prediction and improves the output of  $MLP_{f^{-1}}$ . We further add losses to enforce desired properties of surface parameterization, namely,  $L_{iso}$  provides isometric behaviour,  $L_{angle}$  preserves angles of the faces and  $L_{area}$  preserves face-area (neglecting uniform scaling). The final objective function for surface parameterization is given follows:

$$\mathcal{L}_{uv} = \lambda_1 \mathcal{L}_{cycle} + \lambda_2 \mathcal{L}_{iso} + \lambda_3 \mathcal{L}_{angle} + \lambda_4 \mathcal{L}_{area}. \quad (4)$$

Please refer to our supplementary for description of other loss terms.

## 4 RESULTS & EVALUATION

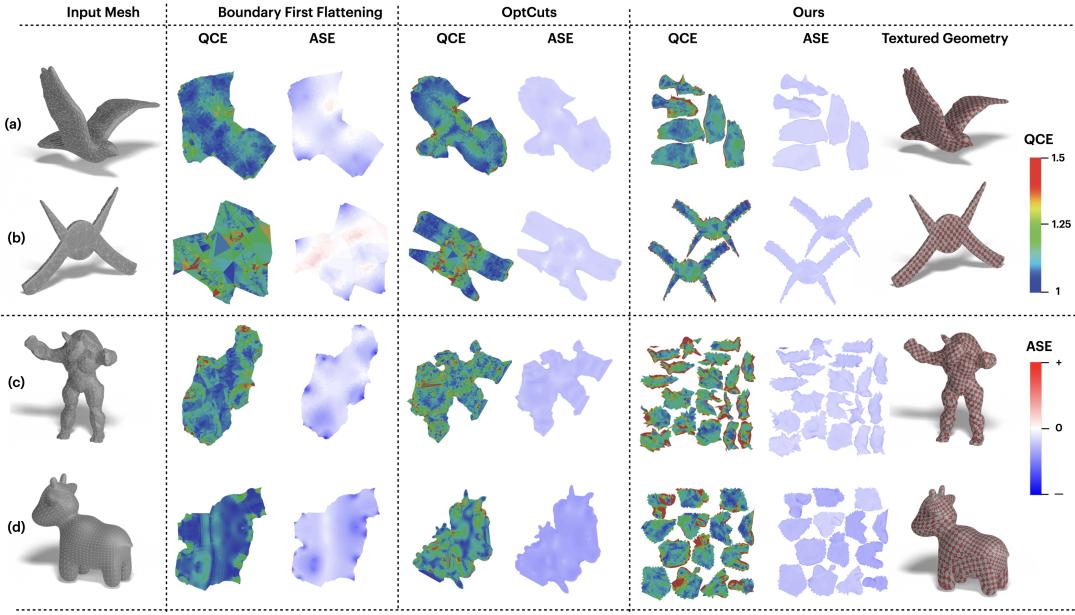
We compute Quasi-Conformal Error (QCE) and Area Scale Error (ASE) on the final texture atlas for quantitative and qualitative evaluation. Please refer to the supplementary for their description.

**Qualitative Comparison:** We compare our framework with BFF[Sawhney and Crane 2017] and OptCuts[Li et al. 2018] in Figure 4. As shown, our framework performs on par with these methods on varying geometrical shapes.

**Quantitative Comparison:** In Table 1, we compare our framework with BFF[Sawhney and Crane 2017] and Opt-Cuts[Li et al. 2018] on a few classes of SHREC [Lian et al. 2011] dataset using QCE and ASE metric. We train our network on 16 meshes for each mentioned class and compute errors on 4 test sample meshes. Please note that, instead of purely object-centric learning, we compare on

Class	BFF		OptCuts		Ours	
	QCE↓	ASE↓	QCE↓	ASE↓	QCE↓	ASE↓
Laptop	1.046	2.052	<b>1.045</b>	<b>2.005</b>	1.196	2.420
Pliers	<b>1.112</b>	1.909	1.128	<b>1.391</b>	1.274	2.895
Rabbit	<b>1.132</b>	2.116	1.160	<b>2.062</b>	1.183	0.992
Scissors	1.156	1.456	<b>1.122</b>	<b>1.276</b>	1.261	2.728
Bird	2.130	<b>1.103</b>	<b>1.129</b>	1.928	1.262	1.996

**Table 1: Comparison of QCE and ASE metrics with BFF [Sawhney and Crane 2017] and OptCuts [Li et al. 2018] on SHREC dataset.**



**Figure 4: Comparison of error plots for QCE and ASE with other methods. First two categories (a) Bird, (b) Pliers are taken from SHREC dataset; (c) Armadillo & (d) Spot.**

a category-specific generalized network, and our performance is comparable to other object-centric methods. Such generalization can be attributed to intrinsic characterization encoded in diffusion features used in our surface parametrization module.

**Discretization-agnostic Learning:** Figure 3 shows the discretization-agnostic learning capability of our framework. We train on a mesh with only  $\sim 3K$  vertices and directly infer at high resolutions ( $\sim 35K$  and  $\sim 100K$  vertices). Please note that the error values for high-resolution meshes stay close to the low-resolution mesh, as observed in the error plots.

More importantly, discretization-agnostic learning allows us to reduce the computation time significantly compared to other methods. Specifically, we train our method on the decimated Stanford’s Armadillo mesh with  $\sim 2K$  vertices and compare our computation time with other two methods at higher resolution ( $\sim 30K$  and  $\sim 100K$ ) as shown in Table 2. *Please refer to our supplementary for implementation details and detailed ablative studies.*

## 5 CONCLUSION

We proposed a novel self-supervised learning based framework for surface parameterization of bounded and unbounded surfaces. Our framework enables discretization-agnostic learning, significantly improving our inference time performance on high-resolution meshes.

**Acknowledgement:** We thank Dhawal Sirikonda for helping us with the visualization of the QCE error metric.

## REFERENCES

- Noam Aigerman, Kunal Gupta, Vladimir G. Kim, Siddhartha Chaudhuri, Jun Saito, and Thibault Groueix. 2022. Neural Jacobian Fields: Learning Intrinsic Mappings of Arbitrary Meshes.  
 Zhiqin Chen, Kangxue Yin, and Sanja Fidler. 2022. AUV-Net: Learning Aligned UV Maps for Texture Transfer and Synthesis.

Resolution	BFF	OptCuts	Ours
30K	17.41 sec	> 10 min	<b>2.92 sec</b>
100K	61.04 sec	> 10 min	<b>5.02 sec</b>

**Table 2: Comparison of computation time for Stanford’s Armadillo.**

- Thibault Groueix, Matthew Fisher, Vladimir G. Kim, Bryan C. Russell, and Mathieu Aubry. 2018. AtlasNet: A Papier-Mâché Approach to Learning 3D Surface Generation.  
 Bruno Lévy, Sylvain Petitjean, Nicolas Ray, and Jérôme Maillot. 2002. Least Squares Conformal Maps for Automatic Texture Atlas Generation. *ACM Trans. Graph.* 21, 3 (jul 2002), 362–371.  
 Minchen Li, Danny M. Kaufman, Vladimir G. Kim, Justin Solomon, and Alla Sheffer. 2018. OptCuts: Joint Optimization of Surface Cuts and Parameterization. *ACM Transactions on Graphics* 37, 6 (2018).  
 Zhouhui Lian, Afzal Godil, Benjamin Bustos, Mohamed Daoudi, Jeroen Hermans, Shun Kawamura, Yukinori Kurita, Guillaume Lavoué, Hien Nguyen, Ryutarou Ohbuchi, Yuki Ohkita, Yuya Ohishi, Fatih Porikli, Martin Reuter, Ivan Sipiran, Dirk Smeets, Paul Suetens, Hedi Tabia, and Dirk Vandermeulen. 2011. SHREC ’11 Track: Shape Retrieval on Non-rigid 3D Watertight Meshes. , 79–88 pages.  
 Pedro V. Sander, John Snyder, Steven J. Gortler, and Hugues Hoppe. 2001. Texture Mapping Progressive Meshes. In *Proceedings of the 28th Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH ’01)*. Association for Computing Machinery, New York, NY, USA, 409–416.  
 Rohan Sawhney and Keenan Crane. 2017. Boundary First Flattening. *ACM Trans. Graph.* 37, 1, Article 5 (dec 2017), 14 pages.  
 Nicholas Sharp, Souhaib Attaiki, Keenan Crane, and Maks Ovsjanikov. 2020. DiffusionNet: Discretization Agnostic Learning on Surfaces.  
 He Wang, Kirill A. Sidorov, Peter Sandilands, and Taku Komura. 2013. Harmonic Parameterization by Electrostatics. *ACM Trans. Graph.* 32, 5, Article 155 (oct 2013), 12 pages.  
 Francis Williams, Teseo Schneider, Claudio Silva, Denis Zorin, Joan Bruna, and Daniele Panozzo. 2018. Deep Geometric Prior for Surface Reconstruction.